Journal of Urban Economics 124 (2021) 103369

Contents lists available at ScienceDirect

Journal of Urban Economics

journal homepage: www.elsevier.com/locate/jue

Why is the rent so darn high? The role of growing demand to live in housing-supply-inelastic cities[☆]

Greg Howard[a], Jack Liebersohn[b,*]

[a] Department of Economics, University of Illinois, 1407 W. Gregory Dr., Urbana, IL 61801, USA
[b] Department of Economics, University of California - Irvine, 3279 Social Science Plaza, Irvine, CA 92617, USA

ARTICLE INFO

JEL classification:
R31
R23
E31

Keywords:
Labor mobility
Housing affordability
Regional inequality

ABSTRACT

Real rents measured in the United States CPI increased 17.4 log-points from 2000 to 2018. We present a spatial equilibrium framework to decompose the increase into several channels, including demand to live in housing-supply-inelastic cities. We find location demand contributed significantly: using parameterizations from the literature and a new rent index, we find it is responsible for between 17 and 73 percent of the overall rent increase, and an even larger share in cities where CPI is measured. The wide range is primarily due to a lack of consensus over the population elasticity to rents, so we estimate it by comparing the effects of demand shocks across cities of differing housing supply elasticities. We find that demand changes have similar effects across cities, suggesting a high population elasticity. Therefore, our preferred estimate is that location demand accounts for more than half of the increase. We discuss implications for housing supply policy.

1. Introduction

The rental price of housing in the United States has risen substantially this century. From 2000 to 2018, the rent component of the consumer price index (CPI) rose 17.4 log-points more than the overall index. Many people's feelings are summarized by the name of a single-issue political party in New York City: "The Rent is Too [Darn] High."

So why is the rent so darn high? Explanations have fallen primarily into two categories. The first explanation is that housing supply was restricted because of increased regulation (e.g., Ganong and Shoag, 2017; Parkhomenko, 2020; Bunten, 2017). The second explanation is a change in the quantity or quality of housing demand, for example demand for houses rather than apartments (e.g. Joint Center for Housing Studies, 2015; Albouy et al., 2016; Gete and Reher, 2018). We propose a third explanation, that the demand for living in housing-supply-inelastic areas has increased.[1] The location demand channel matters because, when a person leaves an elastic city for an inelastic one, the rent in the elastic city falls only a little, but rises more in the inelastic one. So in the aggregate, rents increase. Using a spatial model to decompose the data, we find this location demand channel explains more than half of the national rent increase from 2000 to 2018, and three-quarters of the rent increase in CPI.[2,3]

We first want to establish that demand did rise for housing-supply-inelastic cities. The left panel in Fig. 1 shows that the rise in rents was greatest in cities with an inelastic housing supply.[4] By itself, the fact that

☆ A previous version of this paper circulated under the title "The Geography Channel of House Price Appreciation." Greg would like to thank the National Science Foundation for financial support through the Graduate Research Fellowship (Grant No. 136625). Any opinions, findings, and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the National Science Foundation. All errors are our own. Declarations of interest: none.

* Corresponding author.
E-mail addresses: glhoward@illinois.edu (G. Howard), cjlieber@uci.edu (J. Liebersohn).

[1] The correlation between demand changes and housing supply elasticity has been noted in Davidoff (2016). We are expanding this observation to suggest it can explain a significant increase in the aggregate level of rents.

[2] The CPI measures prices only for certain urban areas, which experienced a relative increase in location demand compared to places not covered by the CPI. We focus on 2000–2018 because real rents increased more quickly then than in previous decades; CPI-rents shows no real rent increase from 1990 to 2000.

[3] As discussed by Molloy (2020), the correlation between local prices and supply restrictions is well-documented. However, our paper considers the causes and implications of such a correlation while taking into account the general equilibrium effects of migration.

[4] To create this graph, we use the definition of Metropolitan Statistical Area (MSA) in Saiz (2010). Using the dataset from his paper, we split all counties into below or above median by the measure of housing supply elasticity that he calculates. We assume that areas outside of MSAs have above-median elasticity. Population data comes from Census estimates. To measure rents, we use the series we create, as described in Appendix B.

https://doi.org/10.1016/j.jue.2021.103369
Received 2 July 2020; Received in revised form 2 June 2021
Available online 16 June 2021

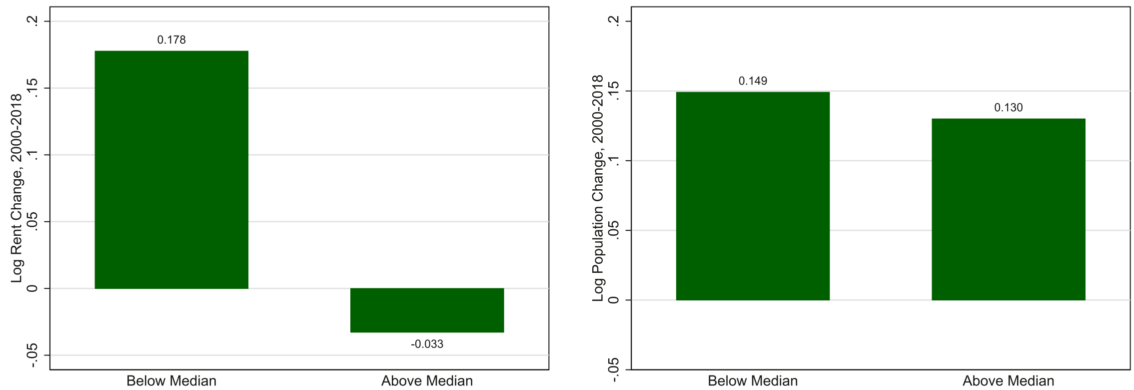

(a) Log Rent Change (b) Log Population Change

Fig. 1. The Change in Rent and Population, by Housing Supply Elasticity, 2000–2018.

rent increases were concentrated in inelastic cities is consistent with any of the three channels. However, we can bring more data to this question. The right panel shows that population also increased more in these cities. This pattern is hard to explain with changing supply restrictions or increases in housing demand, but it is a natural part of our explanation. If people demanded larger houses or if housing became more restricted, we would expect a population decrease in cities where it is harder to build compared to cities where it is easier to build. But the data shows the opposite, which is what we would expect if the demand to live in harder-to-build areas increased.[5]

While Fig. 1 is helpful to make the basic point that demand to live in housing-supply-inelastic areas increased, it does not say how much of the rent increase might be explained by this channel. In fact, our channel could be operative even if populations were growing more in the more housing-supply-elastic regions, precisely because low-elasticity regions are hard to build more housing in.

So we turn to a quantitative general equilibrium spatial model. In our model, individuals make both an intensive-margin choice of how much housing to consume as well as an extensive-margin choice for where they want to live. Housing supply varies across cities, and all individuals must choose to live somewhere. For a set of parameters, the model allows us to decompose the observed changes in rents, housing quantity, and population into four types of shocks: location demand (the demand to live in a particular place), intensive-margin housing demand (what kinds of home each person wants), housing supply, and a national population increase. Including all these shocks allows us to match the changes in population, rents, and housing quantities exactly, and then calculate the moments in the data which tell us why rents changed nationally.

Our model does not have all the sophistications of some urban models because it focuses exclusively on the housing sector. Its strength is in its tractability: we can perform counterfactuals using easily observed moments in the data, similar to a sufficient statistics approach. Critically, we are able to consider the general equilibrium aspect of location demand, that everyone must live somewhere, and maintain closed-form solutions to counterfactuals.

Using the formulae from the model, we learn that the location demand channel is quantitatively important for explaining the rent increase since 2000. For common parameterizations of the key elasticities, our channel explains somewhere between 17 and 73 percent of the increase in rents, and somewhere between 23 and 85% of the increase in rents measured by the CPI. Even at the lower end, this is a substantial proportion of the increase.

However, these intervals are fairly wide, and we wish to be more precise in how responsible the location demand channel is for explaining high rents. We learn from the theory that a critical parameter of our model is the elasticity of population to local rents. This elasticity is defined to be how much, as a percent of the city's initial population, the population changes in response to a one percent increase in rents. Why is this elasticity so important? If it is low, the change in location demand is reflected in populations, so the important moment to quantify the location demand channel is the covariance of population changes and housing supply elasticities, which in the data is small.[6] Under this assumption, an economist would conclude the importance of the location demand channel was at the low end of the range. But if population is highly elastic, location demand changes are reflected in local rents, and the important moment is the covariance of rent increases with housing supply elasticities.[7] Because this covariance is large, an economist would recognize location demand changes as the primary reason rents have increased.[8]

A variety of values for the population elasticity have been used in the literature. Existing parameterizations that we could find range from one-third to infinity,[9] which would lead to very different beliefs on the causes of recent rent increases and the policy implications for ways to lower rents. We also provide estimates of the strength of the location demand channel for the parameters used in these other papers.

Rather than relying on assumptions about the functional form of utility—a widely used approach—we estimate the population elasticity by calculating heterogeneous effects of income changes by housing supply elasticity.[10] This is a strategy similar to Saks (2008) and Diamond (2016), but we study a longer time horizon than Saks (2008) and relax an important assumption in Diamond (2016).[11]

[5] We can also directly measure some things that correspond to this demand increase. In Appendix C, we provide direct evidence that wages and amenities increased more in low-housing-supply-elasticity MSAs.

[6] More precisely, Proposition 1 shows the covariance that matters is population growth with the inverse of the housing supply elasticity plus a constant. This is also small.

[7] The intuition is similar to a supply-and-demand graph. If demand is elastic, then changes in demand are reflected in price even for different supply curves. If demand is inelastic, then changes in demand are reflected in quantities.

[8] It is a model result that a statistic sufficient to calculate the location demand channel are covariances of rents or populations with (transformed) housing supply elasticities.

[9] See Table 1 for a list of papers with different parametrizations.

[10] The total effect on rent to the same change should be proportional to the inverse of the sum of three elasticities: housing supply, intensive housing demand, and location demand. The literature has a much smaller range for intensive housing demand, between 0 and 1.

[11] Saks (2008) estimates the effects of Bartik shocks in cities with high and low housing supply elasticities. She finds that employment responds more in housing supply elastic cities and house prices respond less, both in the short-run. While she does investigate longer-term effects using a VAR structure, the results are not statistically distinguishable.

Table 1
Effect of location demand for Various Parameter Combinations.

Paper	μ	λ	Percent Explained	
			Average	CPI
Suárez Serrato and Zidar (2016), lowest[a]	0.34	1	8	9
Hsieh and Moretti (2019)[b]	1.07	1	17	23
Diamond (2016), College[c]	1.31	1	20	26
Diamond (2016), Non-College[c]	2.50	1	28	38
Suárez Serrato and Zidar (2016), highest[a]	5.1	1	36	50
Caliendo et al. (2017)	∞	1	53	74
Albouy et al. (2016)[d]	∞	2/3	58	77
Van Nieuwerburgh and Weill (2010)	∞	1/2	61	79
Glaeser et al. (2014)[e]	∞	0	73	85

[a] Suárez Serrato and Zidar (2016) provide a large range of estimates in Table 6 of their paper. To calculate the migration elasticity to rent, we take the inverse of their preference dispersion parameter and multiply by the housing share. We report the lowest and highest values from this procedure, corresponding to columns C(5) and B(2), respectively [b] Hsieh and Moretti (2019) report the migration elasticity to wages as $1/0.3$ and the share of housing at 0.32. We calculate their λ as the product of these two. Parkhomenko (2020) also uses this parameterization in his analysis. [c] Diamond (2016) estimates separate elasticities for college and non-college graduates. We put them both here. These are from Column 3 of Table 4. These estimates impose a housing share of consumption, making them more similar to the strategy in Hsieh and Moretti (2019). Without imposing that, the elasticities estimated are slightly larger, closer to 3. [d] Albouy et al. (2016) estimates λ using the assumption of free mobility, i.e. $\mu = \infty$. [e] Glaeser et al. (2014) uses a congestion amenity which plays the same role as μ. They set it to infinity but acknowledge that they are not sure if people like or dislike more population.

If the population elasticity is small, the income change's effect on rent should be larger in a housing supply inelastic area. If the population elasticity is large, there should be no heterogeneity in the rent response. We find a rent increase to income changes that is roughly constant across elasticities. Therefore, our preferred estimate of the eighteen-year population elasticity to rents is infinity, meaning that rents will fully offset any location demand shocks.

Our estimate of this parameter implies that changing location demand is the key reason for the increase in rents since 2000. Overall, we find that more than half of the rent increase is due to this location demand channel, and it was especially important for areas covered by CPI, where it explains three-quarters of the rent increase.

One of the contributions of this theoretical framework is a way to calculate how cross-sectional shocks affect aggregate rents. For example, this method would tell us the contribution of manufacturing to rents via the location demand channel. Typically, economists would start by regressing the local change in rents on some measure of the decline. But to translate that coefficient into a general equilibrium effect requires additional assumptions. In our model, the national population adding-up constraint implies that we also need to know the covariance of the shock with housing supply elasticities.[12]

We also contribute to research on rental prices by developing a new rent index that starts in 2000 and exists for a larger cross-section of cities than is available elsewhere. As noted by Molloy (2020), there is relatively little research on the relationship between housing regulations and rents (as opposed to house prices). We hope that the development of our new index will further the development of this research area.

Our measures are repeat rent indexes we create using multifamily mortgage-backed securities from Trepp. Our series are largely similar to the consumer price index series, but are available for more than 200 cities, rather than the 25 published by the BLS. By measuring rents for a large section of the country, we are able to aggregate local shocks to housing demand and supply in order to estimate their national impact. This would not be possible with more conventional measures. Details of how we construct the series and a comparison to CPI is presented in Appendix B.[13]

While our model was designed to decompose past data, it is also suitable to consider housing-specific counterfactuals. Our estimate of a high population elasticity to rents imply that a local expansion of housing supply is unlikely to have significant long-run effects on rents. Furthermore, expansion of housing supply at the national level has similar effects on rents regardless of where it occurs.[14] Finally, subsidies of housing are more effective at lowering rents when they target housing elastic cities rather than housing inelastic ones.[15]

Our paper is organized as follows: In this next subsection, we review the relevant literature. In Section 2, we present the model that we use to decompose the important forces. In Section 3, we show that the population elasticity to rents is the key parameter to disentangle different causes of the rent increase. In Section 4, we outline our empirical strategy and estimate this elasticity.

Literature Review. The literature has taken a variety of empirical and theoretical approaches towards understanding changes in housing costs. For example, many empirical papers take a cross-sectional approach, seeing how various local shocks, such as regulatory changes or immigration, affect housing prices. More structural papers start with a spatial equilibrium model and estimate the effects of the shock of interest through counterfactual analysis. Our framework focuses on the most critical housing elements of the models, allowing us to make a contribution to both literatures. For the empirical papers, it gives a general equilibrium lens to interpret the coefficients from a cross-sectional regression. For the theoretical papers, it highlights the location preference parameter as central to determining the causes of the rent increase.

Our theoretical contribution builds on the spatial equilibrium framework first developed by Rosen (1979) and Roback (1982). Most closely related to our paper, Van Nieuwerburgh and Weill (2010) develop a spatial equilibrium model showing that wage changes can explain much of the cross-sectional dispersion of house price changes over a long time period. In another related paper, Glaeser et al. (2014) use a spatial equilibrium model to argue that income volatility largely explains house price volatility at the local level. These papers focus primarily on explaining cross-sectional changes.[16] Our paper emphasizes that these cross-sectional changes also have important aggregate effects.[17]

Diamond (2016) explicitly assumes that housing supply elasticities are orthogonal to unobserved local demand shocks. This is a stronger assumption than we need, and it rules out the location demand channel that we have in mind, except to the extent that location demand is captured by the observed shocks she considers.

[12] In the data, it turns out this covariance is significantly positive (Liebersohn, 2017). This means that even though manufacturing might have a negative local effect on rents, it aggregates to a small but positive national effect via the location demand channel. There may also be intensive housing demand changes from manufacturing that this exercise would ignore, but which could still be quantified using our model.

[13] One reason indices similar to the ones we use are not published is that they are subject to measurement error at small levels of geography. And though we do our best to limit it, our indices are also subject to measurement error. One approach we take to limiting the noise is to use empirical Bayes methodology to shrink the estimates based on the uncertainty in the estimation of the rents index. More generally, though, the statistics we consider will be the covariance of rents with housing supply elasticities, and as long as the measurement error is well-behaved, the law of large numbers will mean that our statistics converge. So even if we have insufficient observations to trust our rent index for a specific small city, the measurement error will average out across many cities.

[14] In this case, the location can still have large effects on other outcomes of interest, such as GDP (see Hsieh and Moretti, 2019).

[15] All of these statements use our estimate of population elasticity to rents. However, our theory does provide formulas for calculating the effects of these policies at any value of this elasticity.

[16] The model in Van Nieuwerburgh and Weill (2010) also features an endogenous national equilibrium price, but they strongly emphasize the cross-sectional results of their model.

[17] Another related literature has focused on explaining what causes the income divergence that we have seen over the last forty or so years, mostly arguing that technology is the primary driver with an important role for the sorting of high-skilled workers (Giannone, 2017; Eeckhout et al., 2014). So far, this literature has not focused on the implications for housing costs. Aladangady et al. (2020) document an increase in housing

A related model-based literature uses spatial equilibrium models to estimate how housing restrictions affect cross-sectional and aggregate productivity. Focusing on the cross-section, Ganong and Shoag (2017) argue housing supply regulation can help explain the slowdown in regional convergence. Other papers focus on aggregate productivity or welfare (Herkenhoff et al., 2018; Hsieh and Moretti, 2019; Bunten; Jayamaha, 2019). All four argue that housing restrictions lower aggregate productivity.[18] Like us, they are interested in the aggregate consequences of cross-sectional changes. However, the focus of this literature is on productivity and output, while ours is on housing costs.

Turning to the cross-sectional empirical research, it can largely be split into two categories: estimating local effects of supply and estimating local effects of demand. Starting with supply channels, Saks (2008), Glaeser et al. (2005) and Ihlanfeldt (2007) argue that regulatory changes have made housing production more expensive, raising house prices and causing prices to respond more to housing demand shocks. Diamond et al. (2019) study the effects of rent control on rental prices. Two recent reviews, Gyourko and Molloy (2015) and Glaeser and Gyourko (2018), synthesize the literature on the effects of supply constraints.[19] This is also closely tied to a literature on the relationship between housing supply and housing affordability (Glaeser and Gyourko, 2003; Brueckner, 2009), which is reviewed by Quigley and Rosenthal, 2005 and Molloy (2020).

On the demand side, Gete and Reher (2018) argue that mortgage availability can decrease rents, and Reher (2018) shows that it can raise house prices by increasing the quality of housing. Saiz (2007) shows that immigration shocks can increase rents, and Gorback and Keys (2020) and Li et al. (2021) show that capital inflows raise prices on local housing. Davidoff (2013) and Davidoff (2016) show that many demand-side factors are highly correlated with housing supply elasticities. He focuses on the fact that this ought to rule out using these elasticities as instruments, whereas we are focused on showing that such a correlation has important aggregate effects.

Our theory is flexible enough to accommodate many of the mechanisms studied in these empirical papers, as we would model them as an increase in the demand for housing quantity or a shock to housing supply. In fact, for the supply side, we give an explicit formula for calculating the aggregate effects.

Finally, our paper is related to a large literature on internal migration in the United States. This literature includes both reduced-form and structural research, and studies links between labor markets and migration patterns. Glaeser (2011) gives an overview of the many advantages of moving to cities. In a recent influential lecture, Autor (2019) documents changes in the labor market for college-educated and non-college-educated workers over time. In Appendix C, we show that the location demand shocks we document correlate to observed changes in wages and amenities, consistent with the literature on the determinants of migration (Kennan and Walker, 2011; Molloy et al., 2011).[20]

There is a big literature focused on the reasons for moving. Most related to ours are papers that document that housing costs play a large role in explaining migration patterns (Zabel, 2012; Plantinga et al., 2013). Ferreira et al. (2011) argues that declines in house prices can lock people into their homes. Halket and Vasudev (2014) and Bilal and Rossi-Hansberg (2021) emphasize the life-cycle and investment returns to location. Consistent with that, Nakamura et al. (2021) shows that young people have much larger gains to moving than older people. Monte et al. (2018) argues that commuting openness is an important determinant in the elasticity of employment to a demand shock.

2. Model of housing markets

For our analysis, we use a static general equilibrium spatial model, typical of the spatial economics literature. However, our model abstracts from many features of spatial models, focusing exclusively on the housing sector. The purpose of the model is to be able to decompose the data into several types of shocks, as is commonly done in the business cycle literature to determine the causes of recessions. We can also do counterfactuals by imagining the shocks had been otherwise. One of the contributions that comes from focusing only on the housing sector is that the model is tractable enough so that the counterfactuals can be solved analytically, leading to sufficient statistics for the contribution of location demand.

People in our model, denoted by j, choose a location i in which to live, supply one unit of labor, and consume housing and tradable goods. The price of tradable goods is normalized to 1. Locations differ in their supply curve for housing, their amenity value, and their productivity. In Appendix E, we extend the model to have multiple types of people or multiple types of housing.[21] The utility of an agent is given by

$$U_{ij} = \log u(c_i, b_i h_i, a_i) + \epsilon_{ij}$$

subject to the budget constraint $w_i = c_i + r_i h_i$. The consumption of non-housing goods is c_i,[22] h_i is the consumption of housing, b_i is a city-specific housing demand shifter, a_i is the general amenity level of city i, ϵ_{ij} is a match-specific amenity term, w_i is the productivity of city i, and r_i is the rent in city i. Denote L_i as the number of people that choose to live in city i. We assume ϵ_{ij} is an i.i.d. Gumbel.

This generates a per-capita housing demand function

$$h_i = h_i(w_i, r_i, b_i) \tag{1}$$

which is decreasing in r_i, and an indirect utility function

$$v_i = v_i(w_i, r_i, a_i, b_i) \tag{2}$$

where $U_{ij} = \log v_i + \epsilon_{ij}$. v_i is also decreasing in r_i.

We assume housing is supplied competitively using local land, which is available in a fixed quantity, and the consumption good. Assume the production function is $H_i = Z_i^{\frac{1}{\sigma_i+1}} X_i^{\frac{\sigma_i}{\sigma_i+1}}$, where Z is local land and X is intermediate inputs of the consumption good whose price is already normalized to one. Then from the construction profit first-order condition, the supply curve is

$$\log H_i = \sigma_i \log r_i + \text{constant}_i \tag{3}$$

Note that the housing supply elasticity differs by city.

Local housing markets have to clear, and everyone must live somewhere

$$H_i = L_i h_i \tag{4}$$

$$L = \sum_i L_i \tag{5}$$

inequality over the past fifty years and a lesser increase in rent inequality, but include both inter- and intra-regional inequality.

[18] Bryan and Morten (2019) shows that lowering migration costs can increase aggregate productivity in a developing market context.

[19] More papers studying the cross-sectional price implications of variation in housing supply elasticity include Glaeser and Gyourko (2003), Zabel and Dalton (2011), Glaeser et al. (2012), Guerrieri et al. (2013), Gyourko et al. (2013), Hilber and Vermeulen (2016), and Albouy and Ehrlich (2018).

[20] A subset of this literature has emphasized the heterogeneity within migration patterns. For example, younger and more-educated workers are more likely to move (Chen and Rosenthal, 2008; Molloy et al., 2011). Diamond (2016) shows that different skill groups value amenities differently, Berry and Glaeser (2005) shows that network effects can lead to clustering of high-skilled workers, and Kaplan and Schulhofer-Wohl (2017) discusses location demand for particular skills and occupations. We address concerns regarding heterogeneity in Appendix E.

[21] Other features, such as taxes, can be added as part of productivity or the housing production function.

[22] Here, non-housing goods are all tradable. You could add in non-tradable goods whose equilibrium price depends on w_i or r_i. You would still end up with Eqs. (1) and (2).

Given the Gumbel distribution, the allocation of population to cities is given by:

$$L_i = \frac{(v_i)^\mu}{\sum_j (v_j)^\mu}$$

where $1/\mu$ is the scale parameter of the Gumbel distribution. So

$$\log L_i = \mu \log v_i - \tilde{u} \tag{6}$$

where $\tilde{u}$ is the log of the denominator, which is common across all cities.

Eqs. (1) –(6) define the equilibrium. Given the exogenous wages, i.e. a lack of agglomeration, the equilibrium exists and is unique (Allen et al., 2020).

We take a log-linearized approximation of the indirect utility around the steady-state:

$$\log v_i = \gamma_1 \log w_i - \gamma_2 \log r_i + \gamma_3 \log b_i + \gamma_4 \log a_i + \text{constant}_i,$$

where γs are constants, and its corresponding policy function,

$$\log h_i = c_1 \log w_i - \lambda \log r_i + c_3 \log b_i + \text{constant}_i.$$

where cs are also constants.

These need not be approximations if utility is Cobb-Douglass, in which case $\lambda = \gamma_1 = c_1 = 1$ and γ_2 is the Cobb-Douglass parameter. However, empirical estimates suggest λ closer to two-thirds measured across cities, and the size of a house is subject to large adjustment costs, so imposing Cobb-Douglass may introduce a counterfactually high housing demand elasticity (Albouy et al., 2016).[23] The utility function allows for a normalization, so we normalize $v_2 \equiv 1$.

We are interested in changes in equilibrium, so we express these equations in differences:

$$d \log h_i = -\lambda d \log r_i + \epsilon_i \tag{7}$$

$$d \log H_i = \sigma_i d \log r_i + \xi_i \tag{8}$$

$$d \log L_i = -\mu d \log r_i + \eta_i - d\tilde{u} \tag{9}$$

$$d \log H_i = d \log L_i + d \log h_i \tag{10}$$

$$d \log L = \sum_i L_i d \log L_i = \mathbb{E} d \log L_i \tag{11}$$

where the expectation is initial-population-weighted. These are the five key equations of our model: a housing demand per capita equation, a housing supply equation, a location demand equation, a housing market clearing condition, and a population adding-up constraint. The adding-up constraint is a log-linear approximation.[24,25]

Given data on h_i, L_i, and r_i, there is a unique set of ϵ_i, ξ_i, and $\eta_i - d\tilde{u}$ that can match the data. There are four "shocks" here: η_i is a shock in location demand and is of primary interest. It includes amenity or wage changes. ϵ_i is a shock to housing demand, and includes wage changes and changes in the utility from housing. ξ_i is a local shock to housing supply and includes shocks to available land. The last shock is the change in total population.[26]

These shocks are not the fundamental shocks of the model, but are central to distinguishing the competing hypotheses for why the rent is so darn high. For example, at this point in the paper, we are not interested in whether productivity changes were the fundamental change that caused a rise in rents.[27] We are, however, interested in whether it was a shift in the housing (per capita) demand curves, the housing supply curves, or in location demand that cause the increase in rents. This decomposition will let us attribute the rise in rents to the four channels and allow us to do counterfactual policy analysis.

Our model has abstracted from many features of urban models, including agglomeration, sorting, and housing types. Some of these are easy to map onto our shocks. The model will interpret wage gains from agglomeration as positive location demand and housing demand shocks. A recent literature has also emphasized the endogenous nature of supply (Parkhomenko, 2020). If a location became more desirable which endogenously led to tighter housing supply, the model would interpret that response as a housing supply shock even though the fundamental shock is about location demand. In addition, sorting (e.g., gaining a larger share of college graduates) will be interpreted as a housing demand shock if a city increases its share of people that consumer larger houses. We further discuss sorting in Appendix E. There, we also show how our model accommodates multiple types of housing.

Importantly, the theory allows cities to vary by their housing supply elasticities. This is the key heterogeneity on which we focus. Combining that with the population adding-up constraint will give us our interesting results.

2.1. Model results

Given shocks to η_i, ϵ_i, and ξ_i, we can calculate the effects on rent. Using Eqs. (7)–(10), the change in rents in city i is

$$d \log r_i = \frac{\eta_i + \epsilon_i - \xi_i - d\tilde{u}}{\sigma_i + \mu + \lambda} \tag{12}$$

This formula is going to help us identify μ, the population elasticity to rents. If μ is large, the effects of various shocks on rents will be similar across elasticities, but if μ is small, the effects of a shock will be larger in areas with smaller elasticities.[28] Intuitively, if people are highly willing to move, changes in demand lead to population movement in or out of a city until rent changes offset the wage changes, regardless of the elasticity of housing supply. In the extreme case where $\mu \to \infty$, as in the (Rosen, 1979; Roback, 1982) model, a change in wages should lead to an offsetting change in rents regardless of the housing supply elasticity.[29]

Our primary question is: to what extent is the change in average rent due to shifts in η_i, the location demand channel? To answer this, we invert the model to find the relative η_i's, and consider what the change in rents would have been had all the η_i's been zero.[30] We define the *location demand channel* to be the difference in aggregate rents between the real data and this counterfactual.

To do this with population and rents data, we calculate the η's from the model: $\eta_i = d\log L_i + \mu d \log r_i + d\tilde{u}$. We then change these η's to zero and recompute the equilibrium, in order to quantify the importance of location demand.

[23] Albouy et al. (2016) is estimated across cities. Even in the medium-run, changing the amount of housing consumed in a city requires large adjustment costs, so it may be an overestimate. We also present results for $\lambda = 0$ throughout.

[24] The error induced by the log-linear approximation is approximately half of the variance of population changes, and empirically, it is quite small. In the data, the variance of population growth is a bit over 1 log-point. In all of our counterfactuals, we are considering possibilities in which the variance of population growth is smaller. So given that the average housing supply elasticity is about 2, the overall impact of the approximation error on aggregate rent is at most two-tenths of a log point.

[25] In Appendix E, we consider two extensions to the model and discuss how they would affect our results. In the first, we consider heterogeneity by location elasticity, by including a group of agents that do not move. In the second, we consider a segmented housing market, where some housing is rented and some is owned.

[26] In particular, $\epsilon_i = c_1 \; d \log w_i + c_3 \; d \log b_i$; $\xi_i = d \log Z_i$; and $\eta_i = \mu d \log a_i + \mu v_1 \; d \log w_i + \mu v_3 \; d \log b_i$.

[27] In the last part of the paper we will return to this question and provide evidence on the factors that correlated with changes in location demand.

[28] A high μ does not mean the effects of any shock on rents are small because η_i and $d\tilde{u}$ can scale with μ.

[29] Note that the model-implied η_i's are also changing when μ changes so that the right-hand side of Eq. (12) does not go to zero.

[30] Zero is an unimportant normalization because if all the η_i's increased by 1, then $d\tilde{u}$ would increase by 1 and completely offset it.

We start with two extreme scenarios that simplify the algebra significantly: $\mu = 0$ and $\mu \to \infty$. These correspond to inelastic mobility and perfectly elastic mobility. The formulas for these give us intuition for the important moments in our data. We then generalize our formulas to intermediate values of μ. Proofs are collected in Appendix D.

2.1.1. Population inelastic to rent

proposition 1. If $\mu = 0$, then the contribution of location demand is

$$\text{Location Demand Channel} = Cov\left(d\log L_i, \frac{1}{\sigma_i + \lambda}\right)$$

where the covariance is initial population-weighted.

Here, the effect of location demand is to raise populations more in some regions than others. The increase in population will have a smaller effect on rent in elastic regions as in (12). So if there are bigger population increases in inelastic regions, then location demand will have had a positive effect on average rents.

When $\mu = 0$, most of the interaction between cities is eliminated. η_i determines the population increase directly, because agents are inelastic to the endogenous change in rents. Without the interactions between cities, the model boils down to independent supply and demand curves across cities: Eqs. (7), (8), and (10). Given the excess population, the effect on rents is given by the size of the shock divided by the sum of the demand and supply elasticities: $\frac{1}{\sigma_i+\lambda}(d\log L_i - d\log L)$. Taking the expectation leads to the formula in the proposition.

2.1.2. Population perfectly elastic to rent

proposition 2. If $\mu \to \infty$, then the contribution of location demand is

$$\text{Location Demand Channel} = -\frac{Cov(d\log r_i, \sigma_i)}{\bar{\sigma} + \lambda}$$

where the covariance is initial-population-weighted, and $\bar{\sigma} \equiv \mathbb{E}\sigma_i$ which is also initial-population weighted.

In this case, the key covariance for quantifying the location demand channel is the change in local rents and housing supply elasticity. If rents rise comparatively more in housing inelastic areas, then aggregate rents must rise to make sure enough housing is produced for everyone to have a place to live. This is reflected in the intuition given in the introduction—if one person moves from an elastic to an inelastic place, then rents must rise on average.

This is the case of Rosen (1979) and Roback (1982), with perfect mobility. Any changes in desirability of a place will be reflected in local rents. The housing supply elasticity, along with the local housing demand elasticity, will then determine how much the population will change in response to those rents.

In this situation, which will be our preferred specification later on, the theory provides a key link between cross-sectional relationships and aggregate effects. When $\mu \to \infty$, $d\log r_i$ is the same as $\frac{\eta_i}{\mu}$. So the equation in Proposition 2 provides a framework to consider the effects of a cross-sectional effect on the aggregate. This micro-to-macro approach rests on the national population-adding-up condition, which is the basis for proving the proposition.

2.1.3. Intermediate elasticities

proposition 3. For intermediate values of μ, the contribution of location demand is

$$\text{Location Demand Channel}=\frac{Cov\left(d\log L_i, \frac{1}{\sigma_i+\mu+\lambda}\right) - Cov\left(d\log r_i, \frac{\sigma_i+\lambda}{\sigma_i+\mu+\lambda}\right)}{\mathbb{E}\frac{\sigma_i+\lambda}{\sigma_i+\mu+\lambda}}$$

The limits for this as μ goes to 0 or ∞ correspond to our previous formulas. The formula here is a bit more opaque and in general depends on both the covariance of populations with elasticity as well as the covariance of rents. In the next section, we examine what values the location demand channel takes in the data for various μ's and λ's based on this formula.

2.2. Policy implications

We can use our model to understand the equilibrium effects of counterfactuals. The most policy-relevant questions surround housing supply. In the case of housing supply, the first-order welfare effects for the agents are proportional to the change in average rents, since the effects on the location of workers has only second-order effects by a standard envelope argument. So it suffices to concentrate on the effects on rent. There are also going to be effects on the welfare of landowners, but they will depend on the nature of the housing supply changes, which is outside our model.

A common policy suggestion to make housing more affordable is to expand housing supply. In our model, this corresponds to an increase in ξ_i.[31] Changing ξ_i in one small city has the effect:

$$d\log r_i = -\frac{\xi_i}{\sigma_i + \mu + \lambda}$$

assuming the city is too small to have an effect on $d\tilde{u}$. These effects are decreasing in housing supply elasticity (σ_i), population elasticity (μ) and housing demand elasticity (λ).

We also want to consider a shock to several ξ_i that has an effect on the general equilibrium. In this case, the effect is given by

$$\mathbb{E}d\log r_i = -\frac{\mathbb{E}\frac{\xi_i}{\sigma_i+\lambda+\mu}}{\mathbb{E}\frac{\sigma_i+\lambda}{\sigma_i+\lambda+\mu}}$$

For a finite population elasticity, expanding housing has a larger general equilibrium effect if it expands in low-housing-supply-elasticity places. If population elasticity is very large (i.e. $\mu \to \infty$), this formula reduces to $d\log r = -\frac{\mathbb{E}\xi_i}{\bar{\sigma}+\lambda}$.[32] In that case, the location of the housing supply increase is irrelevant for average rents. That does not mean the equilibrium is unchanged; people will locate in different cities. However, the rent change will be uniform across space and will not depend on where the additional housing is built.[33]

We measure almost all these things in the data. We see $d\log L_i$ in the Census, $d\log r_i$ from our rent series, and $\mathbb{E}d\log h$ from the AHS. We do not have high-quality data on $d\log h_i$. However, we can proxy for it with 2001–2013 changes. As $\mu \to \infty$, the empirically relevant case, the individual $d\log h_i$'s, which we measure poorly, do not matter. In that case all that matters is $\mathbb{E}d\log h_i$ which we measure reasonably well.

A second policy suggestion we consider is the subsidizing of rents. For example, the "Rent Relief Act" (S.3250) introduced by Senator Kamala Harris would subsidize the rents of low-income households above 30% of their incomes. We consider what would happen if rents in some areas are subsidized, financed by lump-sum taxation of the landowners.

---

31 A common result in the urban literature is that changes in regulation actually affect the housing supply elasticity, but are not a shock to the housing supply. If housing supply elasticity changes from σ_i to σ_i', our model will interpret $(\sigma_i' - \sigma_i)d\log r_i$ as a housing supply shock, ξ_i. The propositions hold under this alternate interpretation. However, we will miss second-order effects: for example, a location demand shock combined with a change in the housing supply elasticity.

32 Note that this formula implies that if one city is perfectly housing-supply-elastic and agents are perfectly location-elastic, then housing supply does not affect the average rent.

33 One exercise we will do later is to empirically evaluate the effects of observed housing supply changes over this time period. If we have data on the intensive margin of housing supply, we back out the shocks:

$$\xi_i = d\log H_i - \sigma_i d\log r_i = d\log h_i + d\log L_i - \sigma_i d\log r_i$$

Hence

$$\begin{aligned}\text{Housing Supply Channel} = &-\frac{1}{\mathbb{E}\frac{\sigma_i+\lambda}{\sigma_i+\lambda+\mu}}\Bigg(\mathbb{E}\left[\frac{1}{\sigma_i+\lambda+\mu}\right](\mathbb{E}d\log L + \mathbb{E}d\log h) \\ &+ Cov\left(\frac{1}{\sigma_i+\lambda+\mu}, d\log L_i\right) + Cov\left(\frac{1}{\sigma_i+\lambda+\mu}, d\log h_i\right) \\ &- \mathbb{E}\frac{\sigma_i}{\sigma_i+\lambda+\mu}\mathbb{E}d\log r_i - Cov\left(\frac{\sigma_i}{\sigma_i+\lambda+\mu}, d\log r_i\right)\Bigg) \qquad (13)\end{aligned}$$

In this case, the equilibrium change in rents is

$$\mathbb{E}d\log r_i = -\frac{\mathbb{E}\frac{\tau_i\sigma_i}{\sigma_i+\lambda+\mu}}{\mathbb{E}\frac{\sigma_i+\lambda}{\sigma_i+\lambda+\mu}}$$

where τ_i is the percentage subsidy in city i. In this case, the subsidy times the local elasticity has the identical role to a housing supply increase. Again, if the population elasticity is large, the effect reduces to $-\frac{\mathbb{E}\tau_i\sigma_i}{\bar{\sigma}+\lambda}$. In that case, subsidizing rents in inelastic areas will have a smaller effect on average rents than subsidizing rents in elastic areas. In fact, regardless of the population elasticity, subsidizing rents in elastic areas will be more effective at lowering rents than in places that are inelastic. The reason is that subsidies will lead to more new construction in elastic areas than in inelastic areas, which lowers rents for everyone.

3. The importance of population elasticity

To quantify the location demand channel's contribution to the rent increase of 2000–2018, we apply the formula derived in Proposition 3 to the data. The main takeaway from this section is that the size of the location demand channel will vary greatly depending on the calibration of the population elasticity μ.

3.1. Data

First, we collect data on population growth ($d\log L_i$), housing supply elasticities (σ_i), and rents ($d\log r_i$). We use this to quantify the location demand channel under a variety of assumptions about μ and λ which we take from the literature. Throughout, we use the MSA as our unit of analysis, primarily because the housing supply elasticities at this aggregation have already been estimated by Saiz (2010). Section 4 will show how we estimate μ in the data, which will be the basis of our preferred estimates. Appendix B describes the data sources in greater detail.

We create a new rent index based on the net operating incomes of commercial residential properties. The value-added of our series is that it is available for 217 MSAs, a much greater number than the 25 that the CPI collects.[34] This allows us to calculate how rents covary with elasticity for nearly the entire country, as required by our formulas. The source for this data is CMBS records provided by Trepp. We use an adapted repeat-sales methodology in which we run a regression of log operating incomes on property and year fixed effects in each city, and use the year fixed effects as our city-specific index. In Appendix B, we describe the methodology in greater detail and show that it closely approximates the CPI rents in both the cross section and the time series. Throughout the paper, we present two versions of the rent index. First, we use our unadjusted rent numbers are based solely on the Trepp data. In addition, we improve power in some of our regressions by using an empirical Bayes methodology to shrink them using house price indices. For the majority of cities, this shrinkage makes very little difference, but it does matter for a few outliers that we know are noisily measured.[35] This also allows us to infer rents for the cities where NOI data is not available, so that we can use data from all 269 MSAs where elasticity is available. All of the results are consistent with either set of rent data. If not otherwise specified, statistics in the paper are based on the shrunken rent index. See Appendix B for details.

Housing supply elasticity comes from Saiz (2010). We take these elasticities as given and note that they are estimated using a time sample prior to the one in this paper.[36] Because these elasticities are available by 2000 MSA/Metropolitan Division, this is the basic unit of geography for the empirical estimates. We have to assign non-MSA regions an elasticity because our formulas rely on market-clearing conditions that would not apply if some regions of the country were dropped. We assign them a value of 5.35, which is equal to the 99th percentile of places with a measured elasticity because these regions are primarily small MSAs and non-MSAs with much lower population densities than areas covered by Saiz (2010).[37]

Our results on the importance of location demand do not depend on measuring housing quantity ($d\log h_i$). But for the contribution of housing supply, we need to measure this quantity. We use square footage per person, which we believe is a good proxy for h_i. The important moments for $d\log h_i$ in Eq. (13) are the average national change in square footage per person and the covariance of this quantity with the (transformed) housing supply elasticity by MSA. We measure the former using published tables from the American Housing Survey (AHS) showing median square feet per person, 2001–2017. We calculate the latter ourselves from AHS micro data from 2001 to 2013. Post-2013 AHS public data does not include locations for most housing units, but we believe the 2001–2013 calculations are a good proxy given that the covariance term is close to zero. Appendix B describes the data in detail.

Given the previous theoretical results, the key moments in the data for quantifying the location demand channel will be how rent and population changes relate to housing supply elasticity. To get a sense of the data we plot these relationships in Fig. 2. We show our constructed rent index in the top right, as well as the empirical-Bayes-adjusted rent index in the top left. The patterns are very similar, but there are fewer outliers in the shrunken version of our index. To non-parametrically show the trends, we plot the average change of six quantiles of elasticity.

The increase in rents is concentrated among the least elastic cities. For the least elastic sextile, rents increase by more than 29 log-points. But for each of the three most elastic sextiles, rent actually decreases slightly. For population, the change is hump-shaped, increasing most in areas with elasticity around 2 or 3. The patterns shown in these graphs mean that there is a high covariance between elasticity and rents or prices, but a lower covariance between elasticity and population. This difference will mean that the size of the location demand channel will depend on which covariance is more important for quantifying its effect.[38] A natural question is to what degree does the aggregate population changes hide heterogeneity. In Figure A.1 of Appendix A, we show that the same non-linear pattern holds for a number of subgroups of individuals: immigrants and non-immigrants, college educated and non-college educated, and renters and owners.

One key unmeasured data point is how much rent changes in areas that are outside of MSAs. Given that MSA rent changes are close to zero in the three most elastic quantiles, we will use a rent change of zero for the non-MSA areas as well.[39]

3.2. Estimates

We plug the data into the formula given by Proposition 3 for a variety of values of μ and λ to show how the parameters change the strength

[34] There are not many data sources for rents that cover more than a handful of MSAs back to 2000. The Census reports rents, but it is not quality-adjusted and there are relatively few variables with which to build a hedonic index. Importantly, if we cannot adjust for the quantity of housing as well as the quality, we are likely to measure $r_i \cdot h_i$, the rent times the amount of housing consumed. The repeat rent index avoids that.

[35] The median weight on the Trepp index we create is 87%, and the mean is 84 percent, so for most cities, the shrinkage is minimal.

[36] Other housing supply elasticity estimates are available in Quigley and Raphael (2005), Green et al. (2005), Cosman et al. (2018), Gorback and Keys (2020) and Baum-Snow and Han, 2021. The Saiz (2010) are by far the most ubiquitous in the literature. We check the robustness of our key estimation using the elasticities from Gorback and Keys (2020).

[37] The estimates change very little when varying the value assigned to unmeasured places, for example by assigning them the 90th percentile of MSAs with a measured elasticity.

[38] For the interested reader, there is a statistically significant negative correlation between population changes and housing supply elasticity, consistent with what we showed in Fig. 1. However, the relationship is clearly non-linear, and we think the best way to understand it is with the help of the model.

[39] When we look at house price changes in these areas, they are also close to zero.

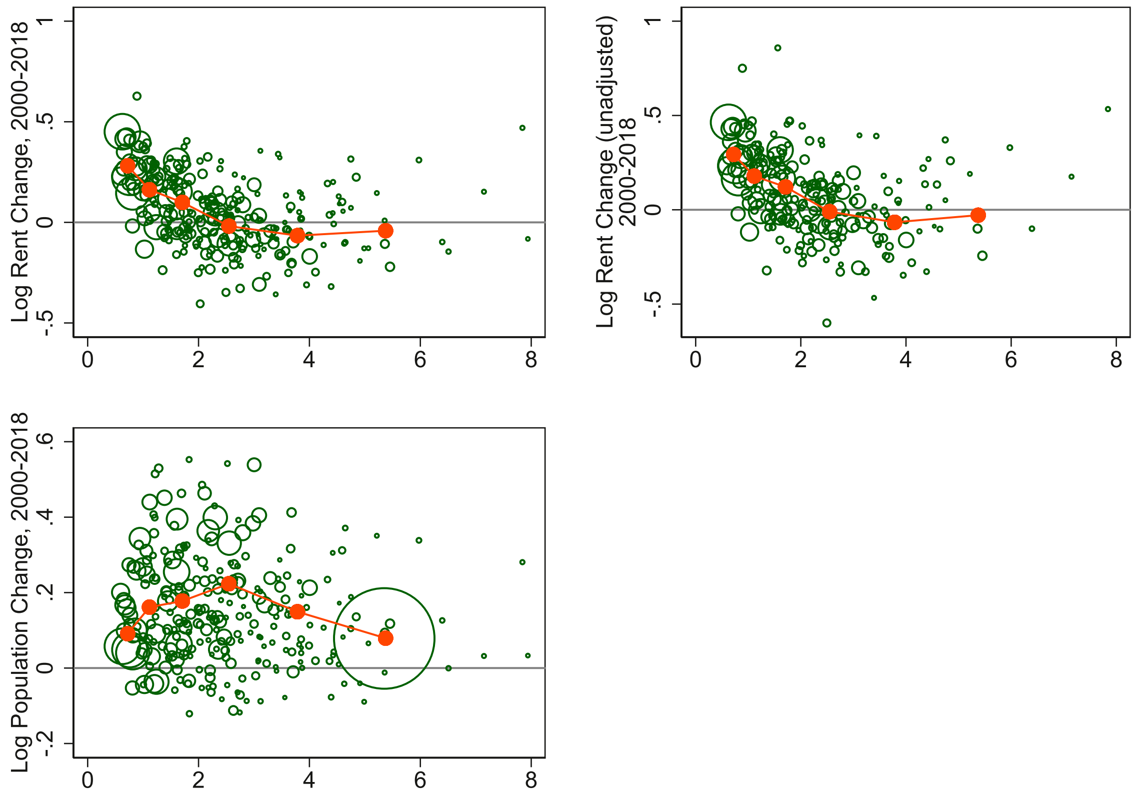

Fig. 2. The relation of rent and population changes (y-axis) with housing supply elasticity from Saiz (2010) (x-axis). The size of the circle represents the MSA's 2000 population. The large circle in the bottom left panel represents all areas outside of MSAs. Averages across sextiles in orange.

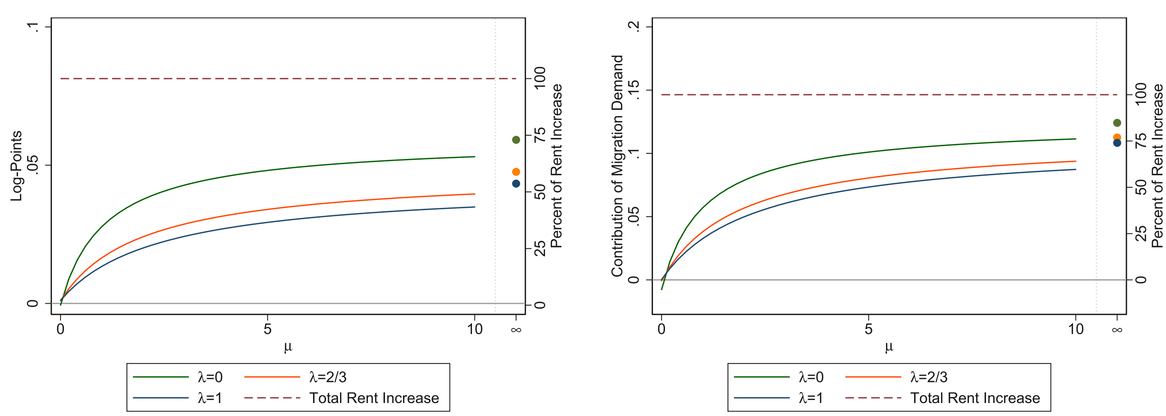

(a) Overall Rents (b) CPI-measured Rents

Fig. 3. The contribution of location demand to the rise in rents.

of the location demand channel. The results are in Fig. 3. The dashed horizontal line at the top of the figure shows the realized total rent increase, and the solid lines show the total rent increase due to the location demand channel as implied by our formulas. We see that the effect is increasing in μ. The dots on the far-right side of the diagram show the rent increase predicted by $\mu = \infty$. Because μ is important for determining the strength of the location demand channel, we provide new estimates for its value in the next section.

λ is the intensive-margin elasticity of housing demand with respect to rents. Albouy et al. (2016) estimates $\lambda = \frac{2}{3}$, which we adopt as our preferred estimate of λ. We also show results for $\lambda = 0$ which would suggest housing per capita is unaffected by rents, and $\lambda = 1$, which corresponds to Cobb-Douglass utility and is common in the literature. As shown in this figure, our results are not that sensitive to the intensive-margin housing demand elasticity. There is little economic difference between our baseline, $\lambda = \frac{2}{3}$, and the largest value used in the literature, $\lambda = 1$. $\lambda = 0$ gives a mildly larger role for location demand.

Given the prominence of rent changes measured in the consumer price index, we also focus on cities in which CPI-rents are measured, in Fig. 3.[40] The overall rise in rents is higher in these cities, but so is the percentage due to location demand, for most values of μ and λ. These

[40] We do not include urban Hawaii and urban Alaska because we do not measure rents in those places.

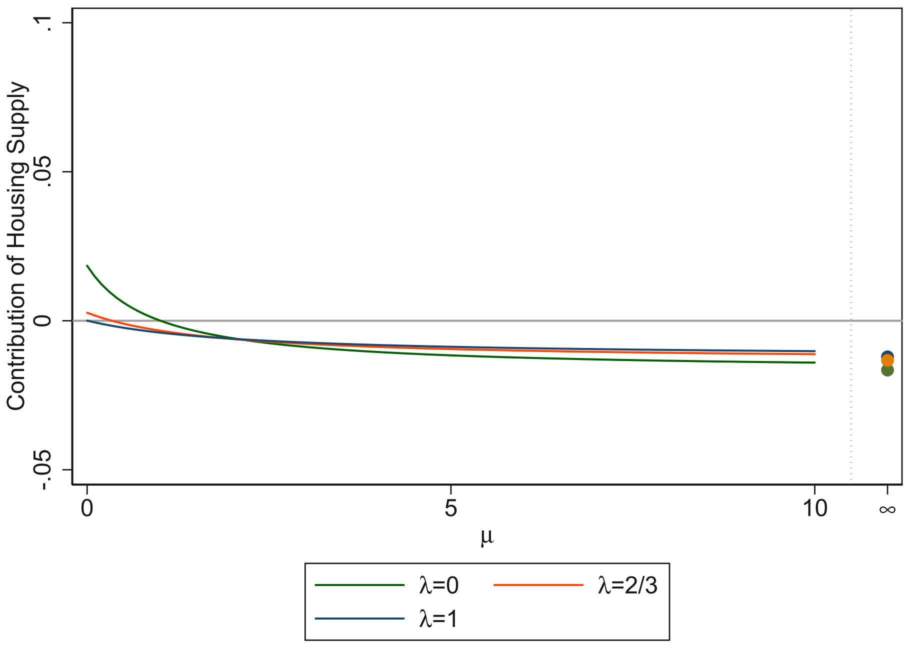

Fig. 4. The contribution of housing supply shifts to the rise in rents.

cities, which are larger and denser than the average city in the U.S., in particular saw some of the biggest increases in location demand.

3.3. The role of housing supply changes

Given our estimates of the location demand channel, the rest of the rent increase can be decomposed into housing supply, intensive-margin housing demand, and population increase. There is particular interest in the literature on housing supply, so we show here how its contribution depends on μ.

The housing supply shocks are close to mean-zero and have a positive correlation with elasticity.[41,42] We apply the formula from (13) to calculate the contribution of these shocks, presented in Fig. 4. As we hinted above, these housing supply restrictions—which are concentrated in *ex ante* inelastic regions—have a larger effect on rents when neither mobility nor construction react very much. When people are fully mobile, the effect is close to zero and actually slightly negative. Overall, however, housing supply shocks play very little role in explaining housing prices. The reason for this is that changes in total housing quantities (that is, population changes times changes in square feet per person) are close to what one would expect given the changes in rents and the housing supply elasticities calculated using historical data (Saiz, 2010).

4. Estimating the population elasticity to rents

From the model, we know that the population elasticity is the key parameter for determining the importance of migration for the rise in rents. There are a large variety of estimates used in existing research, ranging from one-third to infinity. We collect a range of parameters used in the literature in Table 1 to illustrate.[43] For all of these parameters, the location demand channel is quantitatively important. Even using the Hsieh and Moretti (2019) parameters, for which we find the lowest effect of any papers' preferred parameters, the channel accounts for about 20% of the rent increase. Nonetheless, there is still a wide range, and we next estimate the parameter ourselves to improve the precision of our estimate.

The ideal way to estimate the parameter would be to identify exogenous shocks to housing supply, and use that as an instrument for rent changes. Typically, however, there are not large shocks in the data for which the identifying assumptions are plausible.[44] This approach *has* been successfully used to estimate the local, within-city effects of changes to housing supply (Asquith et al., 2021; Mast, 2019; Tricaud, 2021).[45] However, city-wide supply shocks which identify the aggregate population elasticity are harder to identify, which gives rise to alternative approaches used in the literature.

A typical way to calibrate this parameter is to estimate the effects of wage changes on population, and to multiply that number by the share of income spent on housing. The reasoning would be that if rents go up by one dollar, that should have the same marginal effect as if wages went up by one dollar by the envelope theorem. However, this strategy could underestimate the population elasticity to rents because frictions to changing jobs and frictions to changing houses are likely different. For example, a positive shock in one industry may affect demand primarily for the small group of people in that industry, whereas a change in rents likely affects everyone that consumes housing.[46] In practice, approaches which require assumptions about the form of the utility function require both precise measurement of rents and knowledge about all the inputs to utility.

We take a less-restrictive approach, and use Eq. (12) to identify the elasticity. We look for heterogeneous effects of housing demand shocks on rents, by housing supply elasticity. This is a strategy similar to the one

[41] This finding confirms the finding in Ganong and Shoag (2017) that *ex ante* land use regulation correlates with increases in regulation over the course of the 20th century.

[42] When measuring this channel, changes in population are more important than changes in housing quantity per person. The latter are small and overall slightly negative from the years 2001–2017 (median square footage per person decreased from 720 to 700 according to published AHS tables). The covariance between changes in quantity per person and elasticity is close to zero.

[43] A related but distinct literature structurally estimates the mobility response to labor market shocks. Alongside Kennan and Walker (2011), recent research includes Oswald (2019) and Koşar et al. (2021).

[44] We pursued this strategy using local votes to conserve land. While the results were suggestive of a large elasticity, the data was not conclusive.

[45] Anenberg and Kung, 2018 also estimate a high μ within a city, using a more structural methodology. In their setting, a strong preference by particular agents to be in neighborhoods with specific characteristics – including potentially being around neighbors with particular demographics – would imply a low average elasticity of mobility with respect to rents.

[46] A separate concern may be that the housing share of consumption may be hard to measure because it ought to include the effects of housing or other land prices as an intermediate for other consumption. This is one of the reasons there are many estimates in the literature.

pursued by Saks (2008) and Diamond (2016). Eq. (12) states that the effect of a local shock on local rents should be proportional to the inverse of the sum of the elasticity of housing supply, the elasticity of housing demand, and the population elasticity. The heterogeneous effects of a wage shock on rents across housing supply elasticities will give us an estimate of the population elasticity. To put it simply, if a wage shock has the same effect on rents in an elastic and an inelastic city, then the population elasticity must be high. If the wage shock has higher effects on rents in the inelastic city, then the population elasticity must be low. We can derive our regression specification by combining Eq. (12) with the expression in Footnote [26].

Our approach departs somewhat from Saks (2008) and Diamond (2016). We consider a longer time horizon than Saks (2008). We relax the assumption in Diamond (2016) that unobserved labor demand shocks are uncorrelated with housing elasticity.[47] Our assumption is weaker: we require only that changes in unobserved demand shocks are not correlated to the interaction of housing supply elasticity and observed demand shocks. In practice this will mean controlling non-parametrically for housing elasticity in our estimates.

Why not structurally estimate the model instead? Using our approach, we are not using information that we have on housing quantities or population, which could be helpful. There are two reasons not to use housing quantities: first, conditional on the other variables, quantities help us distinguish between local housing demand and supply, but have no information about location demand; and second, we think they are particularly noisily measured (see Appendix B). We will use population, but to do model validation, rather than using it in estimation.[48]

To get a sense of our identifying variation, we run the following regression:

$$d \log r_i = \beta_0 \ d \log w_i + \beta_1 \ d \log w_i(\sigma_i - \bar{\sigma}) + f(\sigma_i) + \nu_i \tag{14}$$

where $d \log w_i$ is a measure of the 2000–2017 change in local wages, conditional on observables, from the American Community Survey,[49] and $f(\cdot)$ is a flexible function so that shocks correlated with elasticity are not influencing the regression. In practice, we control for ten equal-sized bins of elasticity and a linear control.[50] We demean the σ_i on β_1 in order to make β_0 interpretable as the effect of a wage change in an average-elasticity city.

The estimate for how much a wage shock affects the rent is given by $\beta_0 + \beta_1(\sigma_i - \bar{\sigma})$. A negative β_1 would indicate that wage shocks have less effect on rent in elastic areas, indicative that there is a small migration elasticity. A β_1 near zero would indicate that the elasticity is high because the effects are similar in both elastic and inelastic areas.

We present our estimates in Fig. 5. The blue line graphs $\beta_0 + \beta_1(\sigma_i - \bar{\sigma})$, along with 90% confidence intervals. Using either unadjusted rents (right) or using empirical Bayesian shrinkage (left), the main result from the regression is that the blue line is flat; wage changes have similar effects on rents regardless of housing supply elasticity. This would suggest that the migration elasticity is close to infinite. The orange lines present the results from a different specification, in which we split our data into ten evenly-sized bins of elasticity, and within each bin, run a regression of rent changes on income changes. The orange dots show the estimated coefficient, and the lines are 90 percent confidence intervals. These results confirm the previous regression: the effects of income changes seem roughly constant across elasticities.

Given the flat line, a linear approximation to Eq. (12) is reasonable:

$$d \log r_i = \frac{\text{shock}_i}{\lambda + \mu + \sigma_i} \approx \frac{\text{shock}_i}{\lambda + \mu + \bar{\sigma}} - \frac{\text{shock}_i}{(\lambda + \mu + \bar{\sigma})^2}(\sigma_i - \bar{\sigma}) \equiv \beta_0 + \beta_1(\sigma_i - \bar{\sigma})$$

We can then estimate $\mu + \lambda + \bar{\sigma} = -\frac{\beta_0}{\beta_1}$.[51] For a non-negative β_1, our point estimate would be infinity.

Table 2 shows estimates of β_0 and β_1. We check the robustness using two plausibly-exogenous variables that affect local incomes. The decline in manufacturing was a widespread shock that decreased incomes in many regions, and so in columns (2) and (5) we use the manufacturing share as a proxy for the extent of this shock.[52] Columns (3) and (6) use a shift-share (or "Bartik", following Bartik (1991)) shock that predicts local job growth using the interaction of local labor shares and industry-specific wage growth in other regions. The shift-share purges shocks that might affect income and housing costs contemporaneously. Using these variables, we show that rent decreased in manufacturing-exposed areas, and rents increased in areas exposed to growing industries, but that these effects were also not that heterogeneous by elasticity.

Table 3 shows the implied μ's from the regression in Table 2, in the corresponding columns. They are constructed by dividing β_0 by β_1 and subtracting off our preferred value for λ, two-thirds, and the mean elasticity in our data, 2.31.[53] For those on which the sign of β_0 and β_1 is the same, the point-estimate of μ is infinity. In addition, we construct the tenth percentile of estimates via a bootstrap.[54] Drawing samples of our data, we run the same regression and get a distribution of estimated μ's.

Across all six specifications, our point-estimate of μ is infinity or extremely high, suggesting agents are highly mobile in response to changes in rent. Even at the tenth percentile, the estimates of μ are on the higher end of the literature, suggesting highly mobile agents. The exceptions are the results from Columns (4) and (6), where the regression is noisily estimated because we use the non-shrunk rents series. Given these results, our preferred estimate of μ is infinity, meaning that location demand plays a large role in explaining the rise in rents since 2000.

To ensure that the estimates in Table 2 are robustly estimated, Table A.1 in Appendix A shows estimates using more controls and independent variables. Columns 1–2 and 4–5 repeat the estimates in Table 2, adding controls for bins of unavailable land as well as elasticity.[55] Columns 3 and 6 of this table use alternative proxies for income growth, instead of

[47] Diamond (2016) explains that "the exclusion restriction assumes that the level of land-unavailability and land use regulation are uncorrelated with unobserved local productivity changes." Later, she also explains the identifying assumption that "housing supply elasticity characteristics are independent of changes in local exogenous amenities."

[48] Housing supply elasticities from Saiz (2010) are likely noisily measured, and may have shrunk compared to the past (Ganong and Shoag, 2017). Checking the population predictions of the model post-estimation will give us a sense of how large these concerns should be.

[49] See details in Appendix B.

[50] One possible violation of our exclusion restriction would be the effect of agglomeration or congestion, but this effect is likely to be small. In places with higher housing supply elasticity, more people will end up moving in response to an income shock (see Fig. 6). That could generate a correlation between unmeasured shocks to rent (e.g. congestion or agglomeration) and the interaction of the income shock and the elasticity. The reason this would be small is that the amount cities grow in response to a 1% income shock differs only by about 1 percent between the most and least elastic cities, so the agglomeration or congestion change in response to a 1% change in the population would have to be quite large to significantly bias our results. For example, if all the wage-density relationship is the causal effect of agglomeration economies, that would suggest an increase of 1% in population would raise wages by 0.01 percent (Glaeser, 2010), an order of magnitude smaller than our measurement error. While measures for amenities are harder to pin down, we find it unlikely that a 1% change in population has effects large enough to worry us for these regressions.

[51] Note that the scale of the "shock" depends on μ, so if $\mu \to \infty$, β_0 and β_1 do not converge to zero.

[52] The manufacturing shock provides two benefits relative to using the Bartik shock alone. First, the reasons for manufacturing decline are well-studied in the labor economics literature so this shock is more transparent than the Bartik shock. Reasons for the decline in manufacturing employment, which led to lower labor participation in high-manufacturing areas, include import competition with China (Autor et al., 2013) and greater use or robots (Acemoglu and Restrepo, 2020). Second, timing of manufacturing job losses coincided generally with the timing of the 2000s housing boom. The degree to which these two phenomena are interrelated is independently interesting.

[53] To adjust the estimates in Table 3 for a different λ, simply subtract the preferred λ and add two-thirds.

[54] We use 1000 bootstrap samples. The bootstrap is known to be conservative when parameters are at their boundaries.

[55] The unavailable land measure is from Saiz (2010).

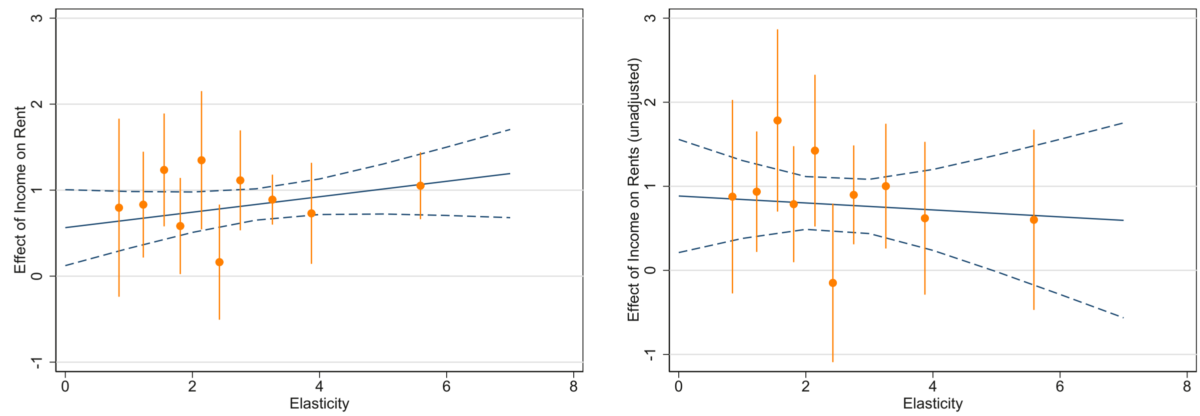

Fig. 5. The effect of wage changes on rent, for different elasticities, with 90 percent confidence intervals.

Table 2
The Heterogeneous Effects of Income Changes on Housing Costs, by Elasticity.

	(1) Rent	(2) Rent	(3) Rent	(4) Raw Rent	(5) Raw Rent	(6) Raw Rent
Income Change	0.792***			0.780***		
	(0.121)			(0.179)		
Income Change by Elasticity	0.0898			-0.0414		
	(0.0767)			(0.150)		
Manufacturing Share		-0.734***			-0.821***	
		(0.158)			(0.232)	
Manufacturing Share by Elasticity		-0.109			-0.229	
		(0.107)			(0.151)	
Bartik Shock			3.440***			4.033***
			(0.624)			(0.989)
Bartik by Elast.			0.742			0.680
			(0.466)			(0.717)
Observations	269	269	269	217	217	217
Controls	X	X	X	X	X	X

Standard errors in parentheses
Robust standard errors. Specifications include controls for deciles of elasticity and a linear control.
* $p < .1$, ** $p < .05$, *** $p < .01$

Table 3
Implied μ's from Heterogeneous Effects Regressions.

	(1)	(2)	(3)	(4)	(5)	(6)
Implied μ	∞	∞	∞	15.6	∞	∞
10th Percentile of $\hat{\mu}$	62	20	84	0	21	0.75

the overall income shock and the manufacturing share. Specifically, we measure income by aggregating a broader range of income-related variables into a single measure using principal components analysis (PCA). Appendix C will use the PCA measure to study the causes of changing locational demand and describes which variables are used to create it. The values of μ implied by these specifications, shown in Table A.2 of Appendix A, are uniformly large. In Columns 7–8, we use the elasticity measure from Gorback and Keys (2020).[56]

Finally, in Columns 9–10 we exploit two components of the Saiz housing supply elasticity—the Wharton Regulatory Land Use Index (WRLURI) and the amount of land unavailable for development—to help deal with measurement error in elasticity. We consider these two variables as noisily-measured proxies for the true elasticity and use one as an instrument for the other to help deal with measurement error.[57]

A major contrast to the previous literature, especially Diamond (2016), is our choice of control variables. Specifically, we flexibly control for housing supply elasticity, whereas such controls are absent from her model. Davidoff (2016) argues that housing supply elasticity is correlated with regional productivity, which could have a direct effect on housing demand. In fact, any channel along these lines is exactly the sort of causal mechanism we are considering in this paper. Therefore, our preferred estimates control flexibly for this elasticity. Appendix Table A.3 estimates μ using specifications that do not have elasticity controls. The estimated μ from these specifications is smaller and more in line with findings from previous papers such as Diamond (2016).

With the previous theoretical results, and our preferred estimate of $\mu = \infty$, we conclude that the location demand channel is responsible for the majority of the increase in rents from 2000 to 2018. For the full set of cities in the sample, we calculate that the location demand channel causes a rent increase of 4.8 log points, which is 58% of the overall rent increase. For the cities where CPI is defined, we find that the location demand channel causes a 10.8% rent increase or 77% of the overall increase in CPI-rents.

The significant role of location demand is robust to the uncertainty from our regression and over the calibration of the housing demand. Recall from Table 1 that if $\lambda = 0$ and $\mu = \infty$, the channel is 6 log points,

[56] This elasticity has few large, negative values, which we set to zero prior to estimation. This measure is available for a much smaller sample of MSAs than the Saiz measure, so we prefer the Saiz measure for our main specifications.

[57] We first scale the WRLURI by projecting the Saiz elasticity onto it. We interact this with income changes in each area to create a version of the elasticity-income shock interaction. Our specification instruments for this interaction uses the interaction of the amount of unavailable land times the Bartik shock.

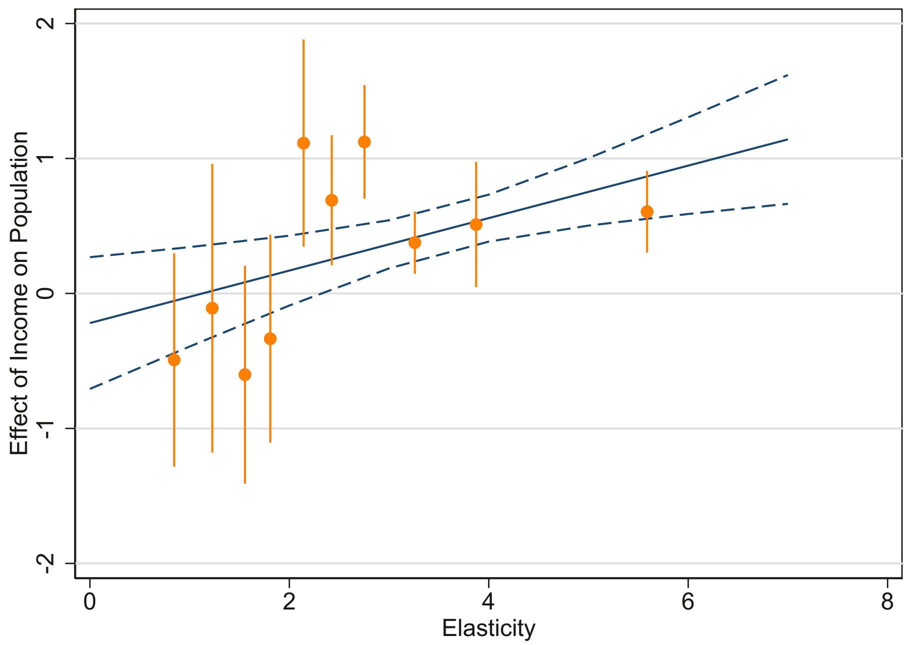

Fig. 6. The effect of wage changes on population, for different elasticities.

and if $\lambda = 1$ and $\mu = \infty$, the channel is 4 log points, which are of similar magnitude to our baseline. We can also directly bootstrap the location demand channel to account for uncertainty in μ: At the tenth percentile of the bootstrap distribution, where $\mu = 62$, the geography channel is 4.6 log points, nearly as large as the median. For cities where CPI-rents are defined, the geography channel is 4.6 log points at the 10th percentile (74%, close to the median value). Even at the 5th percentile, where $\mu = 17$, the geography channel is 4.2 log points, which explains 52% of the rise in rents overall. Using the raw rent measure, which is not shrunk towards house price changes, leads to more uncertainty, but the point-estimate is consistent with our baseline finding. As Figs. 3 and 3 show, these estimates are not affected much by varying the choice of λ. We conclude that regardless of prior beliefs on λ and statistical uncertainty of μ, the location demand channel is explaining a significant chunk of the increase in rents.

4.1. Model validation

Given our estimate of high mobility, our model predicts that population should be more responsive to a location demand shock in more housing-supply-elastic areas. From the model,

$$d \log L_i = (\sigma_i + \lambda) d \log r_i + \xi_i - \epsilon_i$$

We showed that the effect of our shocks on rents was positive and uniform across elasticities. And our identification assumption was that the correlation of our shocks with ξ_i and ϵ_i was the same across elasticities.[58] So the effect of our shocks should be increasing in the housing supply elasticity. We can show this using the same regression as (14). This is shown in Fig. 6. Indeed, we find that the effect of the income shock on population is increasing in housing supply elasticity. While the results qualitatively support our model, the slope of the line in Fig. 6, 0.2, is smaller than our theory predicts, about 0.8. This is primarily driven by the most housing-supply-elastic places; over the least elastic seventy-five percent of cities, the strength of the relationship is consistent with our theoretical prediction. We think there are two possible reasons we see a smaller slope. One is that housing supply might be noisily measured, leading to attenuation bias. Another is that housing supply has become less elastic, a la Ganong and Shoag (2017). We find both of these explanations plausible.

Neither of these possibilities would substantially affect our previous estimation of μ. Comparing the model-implied and empirical slopes provides an estimate of how large either of these concerns should be. The two concerns together lower the slope by about three-quarters. Attenuation in the x-variable of a regression proportionally lowers the regression coefficient. Quadrupling β_1, which is a fairly precise zero, is still zero. Similarly, if the true housing supply elasticity is about a fourth of the real ones, the correct adjustment would be to quadruple β_1, and we would still get a fairly precise zero. Quadrupling would not change our point-estimate for μ at all. The tenth percentile would be closer to 10, which still leads to a very high location demand channel contribution.[59]

In addition to this argument, we also investigate the role of measurement error in biasing our estimate by instrumenting one component of the Saiz (2010) elasticity, the unavailable land, with another, the amount of regulation. This strategy can be found in Appendix Table A.1 and still leads to a high estimate of μ.

Another way of validating the model is to ask whether observable drivers of demand changes match up with the demand shifts implied by the results. In Appendix C, we show that proxies for labor market conditions and amenity growth strongly covary with housing supply elasticity. Then, using the formulas from our model, we quantify the fraction of the location demand channel explained by these factors. We show that a combination of labor-market and amenities changes can explain most of the location demand channel.

Our estimate of μ, while not outside of values used in the literature, is at the upper-extreme and may be surprising to some readers. Importantly, our estimates do not show that everyone has infinite elasticity, only that there is some group that does. As long as such a group exists, the change in location demand will be reflected in rents regardless of elasticity. In Appendix E, we propose an extension of our model in which there are two groups of people, one of which never moves between cities. Nonetheless, when the other group is infinitely elastic, it

[58] Almost certainly, the correlation with ϵ_i, the housing demand shock, is positive, but as long as the correlation is the same for different housing supply elasticities, it should not be a problem. This does mean that we can make predictions on the slope of the relationship, but not the level.

[59] Similarly, measurement error or smaller housing supply elasticities will not greatly affect our estimate of the importance of the location demand channel. Our formula involves the covariance of supply elasticities and rents, so measurement error will add noise, but not bias. We also have housing supply elasticities in the numerator and the denominator, so it is approximately scale invariant.

acts to equalize utility between the cities, and the formula for calculating the location demand channel is exactly the same.[60] Previous literature has emphasized the heterogeneity in moving rates for different groups (Chen and Rosenthal, 2008; Molloy et al., 2011; Diamond, 2016).[61] We can think of μ as estimating the elasticity of the marginal group, which is in equilibrium more elastic than the average person's elasticity.[62]

Because we calculate a high μ, we can plug that into our formulas from before. We conclude that the location demand channel is responsible for more than half of the rent increase, and three-quarters of the rent increase measured by CPI.

5. Conclusions

The past several decades have seen the longest sustained increase in U.S. rents in the post-war period. This paper uses a spatial equilibrium framework of the national housing market to explain why. Our framework allows us to decompose aggregate rent increases into portions explained by changes in housing supply, changes in housing demand, and changes in where people want to live—the location demand channel.

Our key finding is that the location demand channel explains the largest portion of the national increase in rents. Under our preferred parameterization, the location demand channel explains 54% of the rent increase in all U.S. cities and 75% of the rent increase where the CPI rents data exists. These estimates depend on an important unknown parameter: the elasticity of migration with respect to rents. Over the years 2000–2018, we estimate an elasticity of infinity, which corresponds to the parameters in the Rosen-Roback model. We also show that the location demand shocks correspond to increases in wages and amenities in the data.

Our estimate allows us to speak to the policy implications of changing housing supply. First, local expansions of housing supply will have negligible effects on local rents in the long-run. This is because the population elasticity is high, meaning that as more people move in, the marginal person values the housing just as much.[63] Second, national expansions of housing supply can have effects, but it does not matter for rents where the housing is built. A 10% housing supply increase will have the same effect whether it occurs in the most elastic or least elastic regions of the country (of the same population). Finally, subsidizing rent in inelastic cities has small effects on rents paid, whereas subsidies in elastic cities have large effects on rent everywhere. Using the formula from Section 2, a 10% subsidy in the 10% least elastic cities lower aggregate rents by 0.26 log-points, whereas a 10% subsidy in the 10% most elastic parts could reduce rents by 1.63 log-points, more than six times as large an effect.

Our research highlights the importance of the parameter μ for understanding changes in rents. While we think our estimate is an improvement over the existing literature, we hope to see more direct estimates, ideally capitalizing on natural experiments in housing supply.

We would like to thank David Xu for excellent research assistance. We would like to thank Dan Greenwald, Alp Simsek, Ivan Werning, Arnaud Costinot, David Autor, Antoinette Schoar, Jim Poterba, and seminar participants at Ohio State, conference participants at AREUEA, NASMES, MFM, SED, UEA, the UCLA-San Francisco Federal Reserve Housing Conference, the USC-Marshall Junior Macroeconomics Conference, and Midwest Macro for their comments. We would especially like to thank Jessie Handbury, Amir Kermani, and Amine Ouazad for helpful discussions; and Stuart Rosenthal and two anonymous referees for insightful comments.

[60] There is a small adjustment if the two groups consume different-sized housing. However, just the fact that one group consumes smaller housing would not be an issue. It would require that they consume smaller housing and are disproportionately represented in elastic regions. Even then the effect is likely small.

[61] In Figure A.1, we show that on various observable characteristics, the patterns of population changes look fairly similar. Nonetheless, there could very well be heterogeneity within demographic groups.

[62] A related concern is that our model abstracts from the rent-own margin. Also in Appendix E, we present a version of the model with both rental and owner-occupied housing. The lesson from that extension is that we might want to average the user cost of housing with the rents. While the use cost of housing is difficult to measure, Figure B.3b shows that house prices and rents are highly correlated, so we do not think this is a concern.

[63] Each of these policy implications ignores any congestion and agglomeration effects. The literature is not sure of the net sign of these externalities.

Supplementary material

Supplementary material associated with this article can be found, in the online version, at doi:10.1016/j.jue.2021.103369.

References

Acemoglu, D., Restrepo, P., 2020. Robots and jobs: evidence from us labor markets. J. Polit. Econ. 128 (6), 2188–2244.
Aladangady, A., Albouy, D., Zabek, M., 2020. Housing Inequality. NBER Working Paper 21916.
Albouy, D., Ehrlich, G., 2018. Housing productivity and the social cost of land-use restrictions. J. Urban Econ. 107, 101–120.
Albouy, D., Ehrlich, G., Liu, Y., 2016. Housing Demand, Cost-of-Living Inequality, and the Affordability Crisis. NBER Working Paper 22816.
Allen, T., Arkolakis, C., Takahashi, Y., 2020. Universal gravity. J. Polit. Econ. 128 (2), 393–433.
Anenberg, E., Kung, E., 2018. Can more housing supply solve the affordability crisis? evidence from a neighborhood choice model. Reg. Sci. Urban Econ. 80, 103363.
Asquith, B., Mast, E., Reed, D., 2021. Local effects of large new apartment buildings in low-income areas. Rev. Econ. Stat..
Autor, D., 2019. Work of the past, work of the future. Richard T. Ely Lecture given at Annual Meeting of American Economics Association, Atlanta, GA, January, 4.
Autor, D., Dorn, D., Hanson, G.H., 2013. The china syndrome: local labor market effects of import competition in the united states. Am. Econ. Rev. 103 (6), 2121–2168.
Bartik, T.J., 1991. The Debate Over State and Local Economic Development Policies. WE Upjohn Institute for Employment Research.
Berry, C.R., Glaeser, E.L., 2005. The divergence of human capital levels across cities. Pap. Region. Sci. 84 (3), 407–444.
Bilal, A., Rossi-Hansberg, E., 2021. Location as an asset. Econometrica.
Brueckner, J.K., 2009. Government land use interventions: an economic analysis. In: Urban Land Markets. Springer, pp. 3–23.
Bryan, G., Morten, M., 2019. The aggregate productivity effects of internal migration: evidence from indonesia. J. Polit. Econ. 127 (5), 2229–2268.
Baum-Snow, N., Han, L., 2021. The microgeography of housing supply.
Bunten, D., 2017. Is the rent too high? Aggregate implications of local land-use regulation.
Caliendo, L., Parro, F., Rossi-Hansberg, E., Sarte, P.-D., 2017. The impact of regional and sectoral productivity changes on the us economy. Rev. Econ. Stud. 85 (4), 2042–2096.
Chen, Y., Rosenthal, S.S., 2008. Local amenities and life-cycle migration: do people move for jobs or fun? J. Urban Econ. 64 (3), 519–537.
Cosman, J., Davidoff, T., Williams, J., 2018. Housing appreciation and marginal land supply in monocentric cities with topography.
Davidoff, T., 2013. Supply elasticity and the housing cycle of the 2000s. Real Estate Econ. 41 (4), 793–813.
Davidoff, T., 2016. Supply constraints are not valid instrumental variables for home prices because they are correlated with many demand factors. Crit. Financ. Rev. 5 (2), 177–206.
Diamond, R., 2016. The determinants and welfare implications of us workers' diverging location choices by skill: 1980–2000. Am. Econ. Rev. 106 (3), 479–524.
Diamond, R., McQuade, T., Qian, F., 2019. The effects of rent control expansion on tenants, landlords, and inequality: evidence from san francisco. Am. Econ. Rev. 109 (9), 3365–3394.
Eeckhout, J., Pinheiro, R., Schmidheiny, K., 2014. Spatial sorting. J. Polit. Econ. 122 (3), 554–620.
Ferreira, F., Gyourko, J., Tracy, J., 2011. Housing Busts and Household Mobility: an Update. NBER Working Paper 17405.
Ganong, P., Shoag, D., 2017. Why has regional income convergence in the us declined? J. Urban Econ. 102, 76–90.
Gete, P., Reher, M., 2018. Mortgage supply and housing rents. Rev. Financ. Stud. 31 (12), 4884–4911.
Giannone, E., 2017. Skilled-biased technical change and regional convergence.
Glaeser, E., 2011. Triumph of the City. Pan.
Glaeser, E., Gyourko, J., 2018. The economic implications of housing supply. J. Econ. Perspect. 32 (1), 3–30.
Glaeser, E.L., 2010. Introduction to 'agglomeration economics'. In: Agglomeration Economics. University of Chicago Press, pp. 1–14.
Glaeser, E.L., Gottlieb, J.D., Tobio, K., 2012. Housing booms and city centers. Am. Econ. Rev. 102 (3), 127–133.
Glaeser, E.L., Gyourko, J., 2003. The impact of building restrictions on housing affordability. Econ. Policy Rev. 9 (2).
Glaeser, E.L., Gyourko, J., Morales, E., Nathanson, C.G., 2014. Housing dynamics: an urban approach. J. Urban Econ. 81, 45–56.
Glaeser, E.L., Gyourko, J., Saks, R.E., 2005. Why have housing prices gone up? Am. Econ. Rev. 95 (2), 329–333.
Gorback, C.S., Keys, B.J., 2020. Global Capital and Local Assets: House Prices, Quantities, and Elasticities. NBER Working Paper 27370.

Green, R.K., Malpezzi, S., Mayo, S.K., 2005. Metropolitan-specific estimates of the price elasticity of supply of housing, and their sources. Am. Econ. Rev. 95 (2), 334–339.
Guerrieri, V., Hartley, D., Hurst, E., 2013. Endogenous gentrification and housing price dynamics. J. Public Econ. 100, 45–60.
Gyourko, J., Mayer, C., Sinai, T., 2013. Superstar cities. Am. Econ. J. 5 (4), 167–199.
Gyourko, J., Molloy, R., 2015. Regulation and housing supply. In: Handbook of Regional and Urban Economics, 5. Elsevier, pp. 1289–1337.
Halket, J., Vasudev, S., 2014. Saving up or settling down: home ownership over the life cycle. Rev. Econ. Dyn. 17 (2), 345–366.
Herkenhoff, K.F., Ohanian, L.E., Prescott, E.C., 2018. Tarnishing the golden and empire states: land-use restrictions and the US economic slowdown. J. Monet. Econ. 93, 89–109.
Hilber, C.A., Vermeulen, W., 2016. The impact of supply constraints on house prices in england. Econ. J. 126 (591), 358–405.
Hsieh, C.-T., Moretti, E., 2019. Housing constraints and spatial misallocation. Am. Econ. J. 11 (2), 1–39.
Ihlanfeldt, K.R., 2007. The effect of land use regulation on housing and land prices. J. Urban Econ. 61 (3), 420–435.
Joint Center for Housing Studies, 2015. America's Rental Housing: expanding Options for Diverse and Growing Demand. Research Report.
Kaplan, G., Schulhofer-Wohl, S., 2017. Understanding the long-run decline in interstate migration. Int. Econ. Rev. 58 (1), 57–94.
Kennan, J., Walker, J.R., 2011. The effect of expected income on individual migration decisions. Econometrica 79 (1), 211–251.
Koşar, G., Ransom, T., Van der Klaauw, W., 2021. Understanding migration aversion using elicited counterfactual choice probabilities. J. Econ..
Jayamaha, D., 2019. Land-use restrictions: implications for house prices, inequality, and mobility.
Li, Z., Shen, L., Zhang, C., 2021. Local effects of global capital flows: a China shock in the U.S. housing market.
Liebersohn, C., 2017. Housing demand, regional house prices and consumption.
Molloy, R., 2020. The effect of housing supply regulation on housing affordability: a review. Reg. Sci. Urban Econ. 80 (C).
Molloy, R., Smith, C.L., Wozniak, A., 2011. Internal migration in the united states. J. Econ. Perspect. 25 (3), 173–196.
Monte, F., Redding, S.J., Rossi-Hansberg, E., 2018. Commuting, migration, and local employment elasticities. Am. Econ. Rev. 108 (12), 3855–3890.
Nakamura, E., Sigurdsson, J., Steinsson, J., 2021. The gift of moving: intergenerational consequences of a mobility shock. Rev. Econ. Stud .
Oswald, F., 2019. The effect of homeownership on the option value of regional migration. Quant. Econ. 10 (4), 1453–1493.
Mast, E., 2019. The effect of new market-rate housing construction on the low-income housing market.
Parkhomenko, A., 2020. Local causes and aggregate implications of land use regulation.
Plantinga, A.J., Détang-Dessendre, C., Hunt, G.L., Piguet, V., 2013. Housing prices and inter-urban migration. Reg. Sci. Urban Econ. 43 (2), 296–306.
Quigley, J.M., Raphael, S., 2005. Regulation and the high cost of housing in california. Am. Econ. Rev. 95 (2), 323–328.
Reher, M., 2018. The renovation rebalance: how financial intermediaries affect renter housing costs.
Quigley, J.M., Rosenthal, L.A., 2005. The effects of land use regulation on the price of housing: What do we know? What can we learn? Cityscape 69–137.
Roback, J., 1982. Wages, rents, and the quality of life. J. Polit. Econ. 90 (6), 1257–1278.
Rosen, S., 1979. Wage-based indexes of urban quality of life. Curr. Issues Urban Econ. 74–104.
Saiz, A., 2007. Immigration and housing rents in american cities. J. Urban Econ. 61 (2), 345–371.
Saiz, A., 2010. The geographic determinants of housing supply. Q. J. Econ. 125 (3), 1253–1296.
Saks, R.E., 2008. Job creation and housing construction: constraints on metropolitan area employment growth. J. Urban Econ. 64 (1), 178–195.
Suárez Serrato, J.C., Zidar, O., 2016. Who benefits from state corporate tax cuts? A local labor markets approach with heterogeneous firms. Am. Econ. Rev. 106 (9), 2582–2624. doi:10.1257/aer.20141702.
Van Nieuwerburgh, S., Weill, P.-O., 2010. Why has house price dispersion gone up? Rev. Econ. Stud. 77 (4), 1567–1606.
Zabel, J., Dalton, M., 2011. The impact of minimum lot size regulations on house prices in eastern massachusetts. Reg. Sci. Urban Econ. 41 (6), 571–583.
Zabel, J.E., 2012. Migration, housing market, and labor market responses to employment shocks. J. Urban Econ. 72 (2–3), 267–284.
Tricaud, C., 2021. Better alone? evidence on the costs of intermunicipal cooperation.

Further reading

Howard, Greg, Liebersohn, Jack, 2021. Data and code for "Why is the rent so darn high? The role of growing demand to live in housing-supply-inelastic cities". doi:10.3886/E142641V1.